

# Lab 1: Getting Started

For this and all future labs, you will be working on the CS Linux Lab system. If you are not physically in a Linux Lab (FAB 88-09 and 88-10), you may remotely log on to it: `linuxlab.cs.pdx.edu`.

## 1 Setup Your Environment

- Check which shell you are using (bash, csh, ksh, or other):

```
linux> echo $0
/bin/bash
```

This info will be useful for setting up environment variables (see below).

- Check that you have gcc version 7:

```
linux> gcc --version
gcc (Ubuntu 7.3.0-27ubuntu1~18.04) 7.3.0
```

- Check that you have Java version 8:

```
linux> java -version
java version "1.8.0_131"
```

- Check that you have access to the MPI compiler, mpicc:

```
linux> which mpicc
/usr/bin/mpicc
```

- Use addpkg to add the latest version of Chapel compiler to your environment:

```
linux> addpkg
... (a listing of available packages)
```

Use arrow keys and the tab key to select `chapel-1.18.0`; then select <OK> and press enter. To effect the selection, you need to logout your session and re-login.

## 2 Compile and Run Pthreads Programs

To compile Pthreads programs, use gcc with `-pthread` flag:

```
linux> gcc -pthread -g -o sum-pthd sum-pthd.c
```

The compiled programs are run just like regular C programs:

```
linux> ./sum-pthd
```

Note that for this `sum-pthd.c` program, the number of threads is hardwired in the program. In the future, we'll see how to make that adjustable.

**Exercise** Add a `printf` statement in the `worker()` routine to show the thread ID `tid` and the work range `[low, high)` of each individual thread. You may want to place this statement under a control flag:

```
#ifdef DEBUG
    printf(...);
#endif
```

This way, the same program can be compiled to two different versions:

```
linux> gcc -pthread -g -o sum-pthd sum-pthd.c          # regular version
linux> gcc -DDEBUG -pthread -g -o sum-pthd-debug sum-pthd.c # debug version
```

### 3 Compile and Run Java Thread Programs

Compile and run `Sum.java`:

```
linux> javac Sum.java
linux> java Sum
```

#### Exercises

1. This program is not correct. It's missing synchronization. Run this program multiple times until it produces an incorrect result. Can you explain what execution scenario might have caused this incorrect result?
2. Insert a locking statement around the statement `sum += psum`:

```
synchronized (lck) {
    sum += psum;
}
```

The variable `lck` is a synchronization (or monitor) object. In Java, *any* object can be used for this purpose. So we can simply define `lck` as:

```
static Object lck = new Object();
```

Change the program accordingly and re-compile and run.

3. Why do you think the `static` modifier is needed here? Try removing it and re-run the program. What happens?
4. Add a print statement in the `run()` method to print out the thread name and the work range. The thread name can be obtained by calling `Thread.currentThread().getName()`.

### 4 Compile and Run OpenMP Programs

To compile OpenMP programs, use `gcc` with `"-fopenmp"` flag:

```
linux> gcc -fopenmp -g -o sum-omp sum-omp.c
```

Again, the compiled programs are run just like regular C programs:

```
linux> ./sum-omp
```

## Exercises

1. To see the non-intrusive nature of OpenMP, compile the program without the "-fopenmp" flag, and save it in a different target:

```
linux> gcc -g -o sum-omp0 sum-omp.c
```

The result is an identical copy to the sequential version.

One way to confirm that this target code is different from the previous one is to generate and compare their assembly code:

```
linux> gcc -fopenmp -S sum-omp.c      # generate openmp code
linux> gcc -S -o sum-omp0.s sum-omp.c # generate sequential code
linux> wc sum-omp.s sum-omp0.s        # compare their sizes
linux> diff sum-omp.s sum-omp0.s      # compare their contents
```

2. To further confirm that the program sum-omp is indeed running with multiple threads, insert a printf statement inside the for loop to print out the current thread ID, which can be obtained by calling omp\_get\_thread\_num(). For this to work, you also need to include the OpenMP header file:

```
#include <omp.h>
```

*Question:* How many threads are being used?

## 5 Compile and Run MPI Programs

To compile MPI programs, use the command mpicc (which is a gcc wrapper):

```
linux> mpicc -g -o sum-mpi sum-mpi.c
```

Before running MPI programs, you need to setup a host file. Copy linuxhosts to your home directory, and set the following environment variable (different shells use different syntax):

```
linux> export OMPI_MCA_orte_default_hostfile=~/.linuxhosts # bash, ksh
linux> setenv OMPI_MCA_orte_default_hostfile ~/.linuxhosts # csh, tsch
```

You should include this line in your shell startup file to avoid typing it in every time. For bash, ksh, csh, and tsh, respectively, the file names are .bash\_profile, .kshrc, .cshrc, and .tcshrc.

An MPI program is run with the command mpirun, with a flag -n <#copies> indicating the number of copies you'd like to execute:

```
linux> mpirun -n 4 ./sum-mpi // running 4 copies of the program
```

Note that for this program, the number of program copies is specified externally at the time of execution.

**Exercise** Add a printf statement in sum-mpi.c to print out the values of two variables rank (current process id) and size (total number of processes). Compile and run the program to verify that four copies of the code are executed.

## 6 Compile and Run Chapel Programs

Unlike some of the above cases, where gcc handles the compilation, to compile Chapel programs, a separate compiler, chpl, is needed:

```
linux> chpl -g -o sum1 sum1.chpl
linux> chpl -g -o sum2 sum2.chpl
```

For running Chapel programs, you need to set the following env variables in your shell startup file:

```
# for bash, ksh
export GASNET_SPAWNFN=S          # note: no space before and after =
export GASNET_SSH_SERVERS="bevatron boson ..." # list of host names
export SSH_CMD=ssh
# for csh, tcsh
setenv GASNET_SPAWNFN S
setenv GASNET_SSH_SERVERS "bevatron boson ..." # list of host names
setenv SSH_CMD ssh
```

The list of host names need to be manually copied from the file `linuxhosts`.

A Chapel program is run with a flag `-nl <#locales>` indicating the number of locales (*i.e.* hosts) you'd like to use:

```
linux> ./sum1 -nl 1           // running the program over 1 locale
linux> ./sum2 -nl 4           // running the program over 4 locales
```

## Exercises

1. In both programs, the problem domain size `N` is a configurable constant. Try to change it at the time of execution:

```
linux> ./sum1 --N=2000 -nl 1
```

2. Chapel view threads as a lower-level concept, hence does not provide a facility to show which thread a specific code piece is executed by. However, it does provide a facility to show locale information. Add a `writeln` statement inside the `compute()` function in `sum2.chpl` to show where the computation takes place. Use `here.id` to refer to the current locale's name.

## Submission

While labs are not graded, they are required. Everyone needs to submit a lab report. For this lab, you are asked to write a short report (half a page is fine), in plain text or pdf, summarize your experience with these five languages and tools. Which one is your favorite at this point? Also include your answers to the embedded questions in the exercises.

Submit your report through the “Lab1” submission folder on D2L (under the “Activities/Assignments” tab). You should submit your work before the week-end, *i.e.* Sunday 1/13.